

Executive Capability Summary — V4.3

Enterprise AI Incident & Risk Assistant

Executive Overview

Version 4.3 of the Enterprise AI Incident & Risk Assistant has been validated as a governance-first AI decision-support capability designed for use in security-sensitive and regulated environments. The solution is intentionally non-operational and does not automate decisions, recommend actions, or replace human judgment. Its purpose is to reduce decision risk, improve clarity, and strengthen governance.

Validated Capabilities

- Controlled use of information without external assumptions or unstated standards.
- Explicit handling of uncertainty and insufficient evidence.
- Governance-first analysis emphasizing roles, ownership, and accountability.
- Safe incident classification avoiding premature escalation.
- Identification of policy and governance conflicts without autonomous resolution.
- SOC-adjacent operational safety through restraint and audit-ready reasoning.

Risk Controls by Design

V4.3 does not recommend remediation, automate decisions, infer unstated standards or roles, or replace human accountability. These constraints are intentional risk-reducing controls aligned with executive governance expectations.

Executive Suitability

V4.3 is suitable for governance and policy analysis, incident oversight, SOC-adjacent analytical support, and audit, risk, and compliance decision support. Version 4.2 remains the frozen production baseline, while V4.3 is validated for controlled use and consolidation.

Key Executive Takeaway

V4.3 is designed to help leaders and security teams make better, safer decisions by ensuring clarity, restraint, and accountability in complex security situations.